

DIRECTOR OF DIGITAL MARKETING AND SALES

Candidate Assessment

Written responses to five questions, prepared for the interview session the following day.

TRIFLOW
SOLUTIONS

Part A

FIVE WRITTEN RESPONSES

- 01 Proving revenue was caused by marketing
- 02 What can be measured by the eight-week gate
- 03 A lead-generation program that failed
- 04 Engaging a business at partnership
- 05 Adopting technology and what AI cannot replace

Brad Sebastian

Digital Marketing & Business Development · The Woodlands, Texas

713-256-1683

bradsebastian98@gmail.com

01 THE QUESTION

Describe the most rigorous case you have personally made that revenue was caused by marketing — to someone who did not want to believe it. What was your method?

What did it genuinely prove, and what did it not prove? If a skeptical CFO went after that analysis today, where is it weakest — and what would you do differently now?

RESPONSE

The best example I have for this kind of scenario happened recently at NorthSky after we opened up third party channels that generated rapid growth along with new challenges for every department in the company. Attribution was one of those challenges. The rigorous case I made was to create an attribution model in line with a strategy I created to optimize new opportunities created by the third party channels (mostly Amazon). I dubbed the strategy internally "Trojan Horse". The attribution conversation wasn't about assigning credit. It was about informing strategy. So, I made the case before the program ran by working with the sales, customer service, product and accounting teams to change how attribution was assigned to revenue within that strategy (Retroactive attribution is mostly a waste of time unless it's a post-mortem discussion intended to inform strategy).

Our e-commerce revenue had clean attribution where out-of-the box HubSpot models fit perfectly. The Amazon channel was entirely different. Amazon customers were not "our" customers, and we had no ownership or visibility on the individual customer data. We obviously wanted these customers in our channel where the margins were significantly higher. The only way we could capture them was if they managed to find us, which Amazon notoriously makes it very difficult for them to do by design.

IN SHORT

- We changed how revenue got recorded before the program ran, so the question was settled up front.
- A QR code in the box created a customer record where Amazon gave us none.
- The revenue landed in a customer-service pipeline, and the margin was far higher there.

THE THESIS

At NorthSky I changed how the revenue got recorded before the program ran, so every account it created could be counted by anyone.

The Trojan Horse strategy was pretty straight forward, but had a few layers. With the product team we put a QR code in the packaging that led to assembly instructions and warranty registration. Amazon does not

allow sales material in the box, so this was built for service and warranty support, which is what it genuinely was. The registration was a form, so a scan created a customer record where none had existed, and the rest arrived as support tickets.

I created an attribution category called Amazon CX and routed those accounts into a direct pipeline where the margin was far higher. The case had to be made in accounting more than anywhere else. Customer service was a cost center, and there was nowhere to record revenue against a department that was not supposed to produce any. The fix was in the data model: in HubSpot I built an activity type so that a ticket tagged to Amazon was classified as an inside sales activity, which made both the touch and the customer creation attributable.

What it proved was unusually clean. These accounts existed in no system before the program created them, so there was no prior touch history to reconcile; the margin gain per migrated customer was simply the direct-channel margin against the Amazon margin; and the count of accounts created was a report anyone could pull. What it did not prove is what caused the original Amazon purchase; that sale happened inside a channel we could not see. It also did not prove that the migrated customers would not have found us direct on their own. The ones who register a warranty are the engaged ones, and the count says nothing about the people who never scanned.

The sharpest question a CFO could ask today has little to do with the arithmetic. Amazon was by far our largest revenue contributor, and that channel was a deliberate trade of margin for volume, so a program that pulls customers out of it muddies the thing paying most of the bills. The same CFO would also note that an Amazon service function already existed, sitting with our team in China, while this program ran through US customer service, and he would be right to ask whether the same result could have come from the team already closest to the channel, at lower cost.

That is what I would do differently: connect the two siloed service organizations instead of building around them. It was a heavy lift, which is why it did not happen, but it would have put the program where it belonged.

02 THE QUESTION

The business in the data pack has roughly 190 accounts that fit its best-customer profile, a nineteen-month median sales cycle, and a CEO who wants evidence the program is working by a gate review about eight weeks after you start. What do you commit to measuring, and what do you tell him cannot honestly be known yet?

Be specific about the numbers you would put on the page at that gate.

RESPONSE

At eight weeks, projections are hard in a business with a nineteen-month median cycle; 22 of last year's 28 wins had their first contact more than a year before award. Hard is not the same as impossible. A projection can be built from the stages a deal moves through, as long as the stages exist and are defined, and at week eight the honest word for what I would hand finance is a projection, with the assumptions printed next to it, rather than a number.

Here is how I would build it in HubSpot, and it works the same in any CRM. Every opportunity gets a stage with a written definition: first contact, qualified (a definition the salesperson and I agree in writing, since today it is whatever he marked), site visit, quote issued, in procurement, awarded. Each stage carries a close probability that starts as a best guess from last year's record (56 site visits produced 28 awards, so a site visit is roughly a coin flip) and gets replaced by the real rate as quotes and losses are recorded for the first time. Deal amount times stage probability gives a weighted pipeline, and the forecast report rolls it up by expected close quarter. Lead scoring sits on top, ranking the 190 top-tier accounts by engagement (site visits, replies, meetings) so the salesperson works the warmest first and the score itself becomes a leading indicator. At week eight the dashboard shows touches, engaged accounts, MQLs and SQLs by stage, quotes issued, and a weighted pipeline figure with the probabilities beside it, so Alex can question the inputs instead of the output.

IN SHORT

- Projections are hard against a nineteen-month cycle, not impossible: stage-weighted pipeline in HubSpot, with the assumptions visible.
- The week-eight revenue comes from customers who already pay us: 20 handshake conversions and 12 enrollments.
- The rule I care about is at month six, contacted accounts against the ones we held back.

THE THESIS

At week eight I hand finance a projection built from pipeline stages, with the probabilities printed beside it, and the number itself at month eighteen.

The revenue I can actually commit to comes from customers who already buy from us, where the cycle does not apply: 480 restored assets at active customers with no maintenance, 176 assets on a handshake with no contract and no escalator, and 41 agreements coming due with no renewal process. Maintenance recurs and earns 64% gross margin against 27% on Restoration, so it is also the best revenue in the business.

The eight-week figures are small on purpose, because Exhibit 4 puts the median wait from restoration to maintenance at fourteen months. On the page: 20 of the 176 handshake assets moved to a signed agreement, since those customers already pay every year and the ask is a signature; 12 of the roughly 240 wave-one assets enrolled or in procurement; all 41 renewals contacted, with an expiry date on all 248 agreement records; the first 26 of the 52 lapsed customers called, with the other half held for the month-six comparison. Next to each number goes the rule. Fewer than 10 handshake conversions means the offer is wrong, probably the escalator. Under 6 enrollments means I pause the first wave and figure out why, rather than kill it against a fourteen-month median.

The plumbing that makes that dashboard trustworthy goes in first: an owner and a contact on every one of the 480 assets, 176 handshake assets, 41 renewals and 52 lapsed customers, source as a required field on every new record (only 61% of wins have one today), every quote written down, losses kept instead of deleted, and a named owner with a promised response time on every inbound inquiry.

What I would not put a number on at week eight is marketing-attributed new-customer revenue or a cost per acquisition. Those need the quotes and losses nobody has kept and a cycle that has finished, so the first version I would put my name to comes around month eighteen. The projection gets him there honestly in the meantime.

The number I care most about comes at month six. Half of the accounts holding those 480 assets get the Maintenance offer in Q1 and half wait, and at month six I compare how many assets each half enrolled. Eight points or more ahead and the cross-sell gets the budget. Under three, the budget goes somewhere else, and I would say so in the same meeting.

03 THE QUESTION

Tell us about a lead-generation program you built that failed. How did you know?

How long did it take you to know? What had it cost by then? What did you kill, change, or keep, and who did you have to convince?

RESPONSE

When I joined NorthSky the company was a startup with no customers and no sales history, and the plan was to reach as many buyers as we possibly could, as fast as we could. We had just invested in ZoomInfo, the sales director set the filters, and tens of thousands of records came into HubSpot. I built the sequences and started sending. The plan was the company's; the execution was mine, and I own how it started.

The first results told the story. Open rates were very low, clicks were lower, and conversions were close to nothing. Cold records at that volume, with no engagement history behind them, were never going to perform, and it did not take long to see it. I knew inside the first reporting cycle, which is the advantage of running everything through one CRM: the numbers were in front of me the same week the sends went out.

IN SHORT

- A large cold import into a new domain, and the first results were poor.
- I changed the logic fast: engagement-based workflows, hygiene filters, and Warmy for deliverability.
- The ZoomInfo data still informed sales. The engaged campaigns are what pulled the numbers up.

THE THESIS

The cold sends did not perform. I saw it in the first reporting cycle and moved the program to people who had already engaged.

The cost was modest, which is the part I am glad about. A short run of sends that did not perform, and some early drag on deliverability before I got ahead of it. It never became a reputation problem, because I moved on it fast.

What I killed was the volume. What I changed was the logic. I built HubSpot workflows that recorded engagement on every contact and dropped anyone who was not engaging, so the unusable records could not weigh down the ones worth talking to, and I ran the list through hygiene filters before anything else went out. On deliverability I used Warmy and a couple of other tools to warm the sending domain properly and keep it healthy. What I kept was the ZoomInfo data itself. It still informed sales, as account intelligence for the people who had shown interest, rather than as a list to blast, so the investment was not lost.

Then I pointed the program at people who had done something: abandoned carts, abandoned product pages, site visitors, anyone who had opened or clicked. Those campaigns are what pulled the numbers up, from about 5% opens to just under 22%, and clicks from about 0.2% to nearly 5%.

The people I had to bring along were the sales director and the leadership team, who had paid for the data and wanted the reach. What got them there was not a conversation about deliverability. I showed them that a much smaller list of people who had engaged was producing more conversations than the cold sends ever did, and that the reorder campaigns built on the same logic were converting at 10 to 12%. Fewer emails, more revenue.

04 THE QUESTION

The companies we acquire and partner with will be at various maturity stages, have a variety of tools (or lack of tools), and a wide mix of structures and personalities. How do you go about engaging with and understanding the businesses when we initially partner with them?

RESPONSE

I have been on the inside of this twice: once at Integrated Global Services, where we were the platform company absorbing two businesses that had grown up separately, and once at NorthSky, where three house brands were launched together on one engine.

At IGS we absorbed Cetek and then Hot-tek, which made three companies selling overlapping services to the same refineries, often to the same maintenance managers, sometimes on the same shutdown. I ran the rebrand that folded Hot-tek into IGS and the marketing side of bringing three operations into one commercial motion. At NorthSky the three brands ran through one demand engine, one CRM and one reporting standard instead of three parallel programs.

IN SHORT

- I start with the people who hold the relationships, before I open a system.
- I inventory the tools, data and reports before changing any of them; at STI the valuable thing is twenty years of performance data.
- Across companies I hold the CRM, the qualified-lead definition, the source field and the reporting standard fixed. The rest can flex.

THE THESIS

I ask the people who hold the relationships to walk me through their last ten wins and ten losses before I open a single system.

What I took from both is that the first month is for understanding, and the changes come later. I start with the people who hold the relationships, because in a company like STI the president is the CRM and the salesperson is the pipeline. I ask them to walk me through their last ten wins and their last ten losses in their own words before I open a single system, which tells me how the business actually sells, who refers work in, and where the handoffs break. It also tells me what they are proud of, which matters more than it sounds.

Then I inventory what exists: tools, data, collateral, the website, the report somebody runs every Monday. I am not there to rip any of it out yet. I want to see it, because in every one of these businesses there is something valuable that nobody has packaged. At STI it is twenty years of before-and-after performance data on more than a thousand assets, and the byline of a technical authority that still ranks in search. The fastest way to earn the right to change things is to fix something they feel first, like the form that does not resolve or the renewal nobody knew was coming.

The platform pieces come after that. There are only a few I insist on across every company: one CRM, one definition of a qualified lead, and the same source field and reporting standard everywhere, because the day TriFlow buys its third company somebody will want to compare it with the first two on one page. Everything else I am happy to leave alone for a while. If a brand name means something to a cautious buyer, it stays. If a tool works and the team already uses it, it stays until there is a real reason to move it. What the Hot-tek rebrand at IGS taught me is that customers need to hear why the name is changing before they see the new one, and I would take that slowly with a refinery maintenance manager who has bought from the old name for years.

05 THE QUESTION

AI tools have changed how we approach almost everything, and they have real downsides as well. Describe your approach to adopting new technology. And separately: what do you believe you bring to the table that cannot be replicated by a good AI prompt?

RESPONSE

I learn a tool by running it on my own work before it goes anywhere near a customer, and I write the rules down as I go, usually after it has failed in some way I did not expect. The best example is the one I am living in right now. Since July I have built and run a system of AI agents that manages my own pipeline from end to end, and I built it bigger than the job needed so I could see where it breaks. There is one record of truth for every opportunity, the documents are generated from code so a fact corrected once is corrected everywhere, the dashboard rebuilds itself from the data, and the agents follow written procedures for research, contact vetting, drafting and rehearsal. The agents can reach my email, cloud storage, a browser and my local files through Model Context Protocol, under a permission model I designed, and nothing goes out to another person until I have read it and said go.

What I would bring to TriFlow from all that is the rules, more than the plumbing. The mistake I see most often is the model rewording a true fact into something nearby that sounds like what the reader wants to hear. Telling it to be careful does nothing. What works is a list of approved names and terms and a check that runs every time, whether or not anyone remembers to look. A tool name I had never used made it into twelve documents before I caught it, and the rule that came out of that is that only I can add a tool to the verified list. I also test everything where it will actually be used before it ships, because something that works where it was made and breaks the moment it lands is a failure.

IN SHORT

- I built and operate a multi-agent system with an approval gate on every outbound action.
- A tool name I had never used got into twelve documents before I caught it; now only I can add a tool to the verified list.
- Every substantive correction in my system has come from me knowing my own record better than the model does.

THE THESIS

Since July I have run my own pipeline on a system of AI agents, and nothing leaves it until I have read it. Every rule in it came from a mistake it made.

None of the rulebook was written in advance. Every rule in it is a mistake the system made once, written down so it cannot make it twice. The agent drafts, checks its own work and sends under my approval, and it gets a little more room each time the record shows it has earned it. That is how I would use agents at STI. They would research the named accounts in the plan and personalize the outreach, a person would read everything before it went out, and every time the system got something wrong the fix would go into the rulebook.

What I bring that a prompt cannot, first of all, is the facts. Every real correction in that system has come from me knowing my own story better than the model does. Then it is the judgment about what to aim at. My system is tuned to earn meetings, because my own evidence says meetings are what convert, and no prompt made that call. I also carry the consequences of whatever goes out the door. And I can tell somebody no, which a model will not do, and deciding what the business will not do is a person's job.

Note on AI use

Claude is my primary AI engine. Over the past few months I have built an efficient work engine inside it, with its own rules, skills and gates, that makes everything I produce faster and more accurate. I applied my drafts on this assessment to that existing framework. I produced this document; the engine and its agents enhanced it.

ENGINE, GATES AND PASSES APPLIED TO THIS ASSESSMENT

STAGE	WHAT IT DOES	TOOLS, RULES AND GATES APPLIED
Sourcing	Loads the assessment and its exhibits as the only working material.	Clean-room handoff Verbatim source transcription Single-résumé claim pool
Data and arithmetic	Builds the model from cited inputs and recalculates on every change.	Cited-inputs workbook Assumption labels A1 to A16 Recalculation check: 233 formulas, zero errors
Drafting	Turns my notes and dictation into drafts inside the assessment's own constraints.	Rubric-to-outline map Word-guidance rule Punctuation and vocabulary rules
Verification	Rebuilds and checks every output before anyone reads it.	Verify gate: page limit, word count, punctuation, vocabulary, audience Figure-to-workbook match
Review	Reads the finished draft the way the recipient would.	Cold-read reviewer agent Rubric scoring, two independent scorers Numbers escalation rule
Proofing cycles	Applies the findings and repeats until the review comes back clean.	Guarded editor agent: no new facts, no number changes Monte Carlo document check: three rounds plus a scoring pass
Design	Layout and typography.	Claude Design

A catalog of the rules, skills, gates, operations and workflows I have built inside Claude can be reviewed at bradsebastian.com/agentic.

Brad Sebastian

Revenue, engineered. · bradsebastian.com

END OF PART A